

Concurrent Programming III - Differing Tasks

HPPS

David Marchant

Approaches to concurrency

- So far we've looked at OpenMP and threading
- These are what we use most often, especially where we simply want to go faster

Reminder of problems so far. . .

- **Classical problem classes of concurrent programs:**
 - **Races:** outcome depends on arbitrary scheduling decisions elsewhere in the system
 - Example: who gets the last seat on the airplane?
 - **Deadlock:** improper resource allocation prevents forward progress
 - Example: traffic gridlock
 - **Livelock / Starvation / Fairness:** external events and/or system scheduling decisions can prevent sub-task progress
 - Example: people always jump in front of you in line
- **As well as these, some problems just aren't suited to identical threads, we want different processing doing different things.**

Threads vs. Processes

■ How threads and processes are similar

- Each has its own logical control flow
- Each can run concurrently with others (possibly on different cores)
- Each is context switched

■ How threads and processes are different

- Threads share all code and data (except local stacks)
 - Processes (typically) do not
- Threads are somewhat less expensive than processes
 - Process control (creating and reaping) twice as expensive as thread control
 - Linux numbers:
 - } ~20K cycles to create and reap a process
 - } ~10K cycles (or less) to create and reap a thread
 - } *Much* larger difference on non-Unices.

Threading example

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

int THREADS = 4;

struct thread_args {
    pthread_t thread_id;
    int thread_num;
};

void *my_thread(void *arg) {
    struct thread_args *thread_info = arg;
    printf("Hello from thread %d\n", thread_info->thread_num);
    return NULL;
}

int main() {
    struct thread_args *all_thread_info = calloc(THREADS, sizeof(*all_thread_info));
    pthread_t* thread_nums[THREADS];
    for (int i=0; i<THREADS; i++) {
        all_thread_info[i].thread_num = i + 1;
        pthread_create(&all_thread_info[i].thread_id, NULL, my_thread, &all_thread_info[i]);
    }

    for (int i=0; i<THREADS; i++) {
        pthread_join(all_thread_info[i].thread_id, NULL);
    }

    free(all_thread_info);
    exit(0);
}
```

```
user@system:~ gcc -o example example.c -lpthread
user@system:~ ./example
Hello from thread 2
Hello from thread 1
Hello from thread 3
Hello from thread 4
user@system:~
```

example.c

pthread vs. OpenMP

- Generally OpenMP is:
 - Quicker and easier to implement
 - Good enough
 - Handles a lot of races and deadlocks for you
- Manual pthreads is:
 - Not actually that hard
 - A lot more verbose
 - Lets you make all manner of errors
 - More scope for fine tuning
 - You aren't just limited to parallelising loops

but OpenMP isn't as limited as I first suggested

What we've seen so far

```
void vector_add(int n, const int *a, const int *b, int *c) {  
    #pragma omp parallel for  
    for (int i = 0; i < n; i++) {  
        c[i] = a[i] + b[i];  
    }  
}
```

- This function adds two vectors A and B into C
- For N items we start N parallel regions
- Possibly each region is one thread, but this is hardware/scheduler dependent

Dividing Into Chunks

```
void vector_add_quartered(int n, const int *a, const int *b, int *c) {  
    int chunks = 4;  
    int chunk_size = n / chunks;  
    #pragma omp parallel for  
    for (int chunk=0; chunk<chunks; chunk++) {  
        int start = chunk * chunk_size;  
        int end = start + chunk_size;  
        for (int i = start; i < end; i++) {  
            c[i] = a[i] + b[i];  
        }  
    }  
}
```

- Here we divide the problem N into 4 chunks
- Assuming we've 4 cores, this could neatly map with the minimum threading overhead

Programmatic Chunks

```
void vector_add_dynamic(int n, const int *a, const int *b, int *c) {  
    #pragma omp parallel {  
        int t = omp_get_thread_num();  
        int P = omp_get_num_threads();  
        int chunk_size = n / P;  
        int start = t * chunk_size;  
        int end = start + chunk_size;  
        if (t == omp_get_num_threads()-1) {  
            end = n;  
        }  
        for (int i = start; i < end; i++) {  
            c[i] = a[i] + b[i];  
        }  
    }  
}
```

- Better still would be to dynamically detect how many chunks to split into
- Note we need to account for an uneven divide in our last chunk
- Also no longer using **for** clause

Why bother?

- Avoids the overhead of starting and stopping loads of unnecessary threads
- Can be used to gives different problems to different threads
- Usually we want the same *task*, but on different *data*
- This example could be achieved with a reduction clause, but we could use the same logic to manage problems that reductions can't

Key functions

```
int omp_get_thread_num(void);
```

- Gets a unique id number for the calling openmp thread
- Must be called within openmp scope

```
int omp_get_num_threads(void);
```

- Gets count of openmp threads in current openmp scope
- Must be called within openmp scope

```
int omp_get_max_threads(void);
```

- Gets count of theoretical max scheduled by openmp in the current program
- Can be called anywhere

Note that all of these require:

```
#include <omp.h>
```

Threads vs. Processes

■ How threads and processes are similar

- Each has its own logical control flow
- Each can run concurrently with others (possibly on different cores)
- Each is context switched

■ How threads and processes are different

- Threads share all code and data (except local stacks)
 - Processes (typically) do not
- Threads are somewhat less expensive than processes
 - Process control (creating and reaping) twice as expensive as thread control
 - Linux numbers:
 - } ~20K cycles to create and reap a process
 - } ~10K cycles (or less) to create and reap a thread
 - } *Much* larger difference on non-Unices.

Processes in C

- **Relatively few important functions**
 - `fork()` - creates a duplicate child process
 - `execve()` - replace current process with specified program
 - `getpid()` - get current process id
 - `wait()` - Wait for any child processes to complete
 - `waitpid()` - Wait for any or specific child process to complete
- **Note that these won't work on windows, so we won't really focus on them**
- **Even on non-windows, you're probably better using threads as they are much faster**
- **Processes tend to be easier to program though (ish)**

Processes in C

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    fork();
    fork();
    printf("Hello from process %d\n", getpid());
    exit(0);
}
```

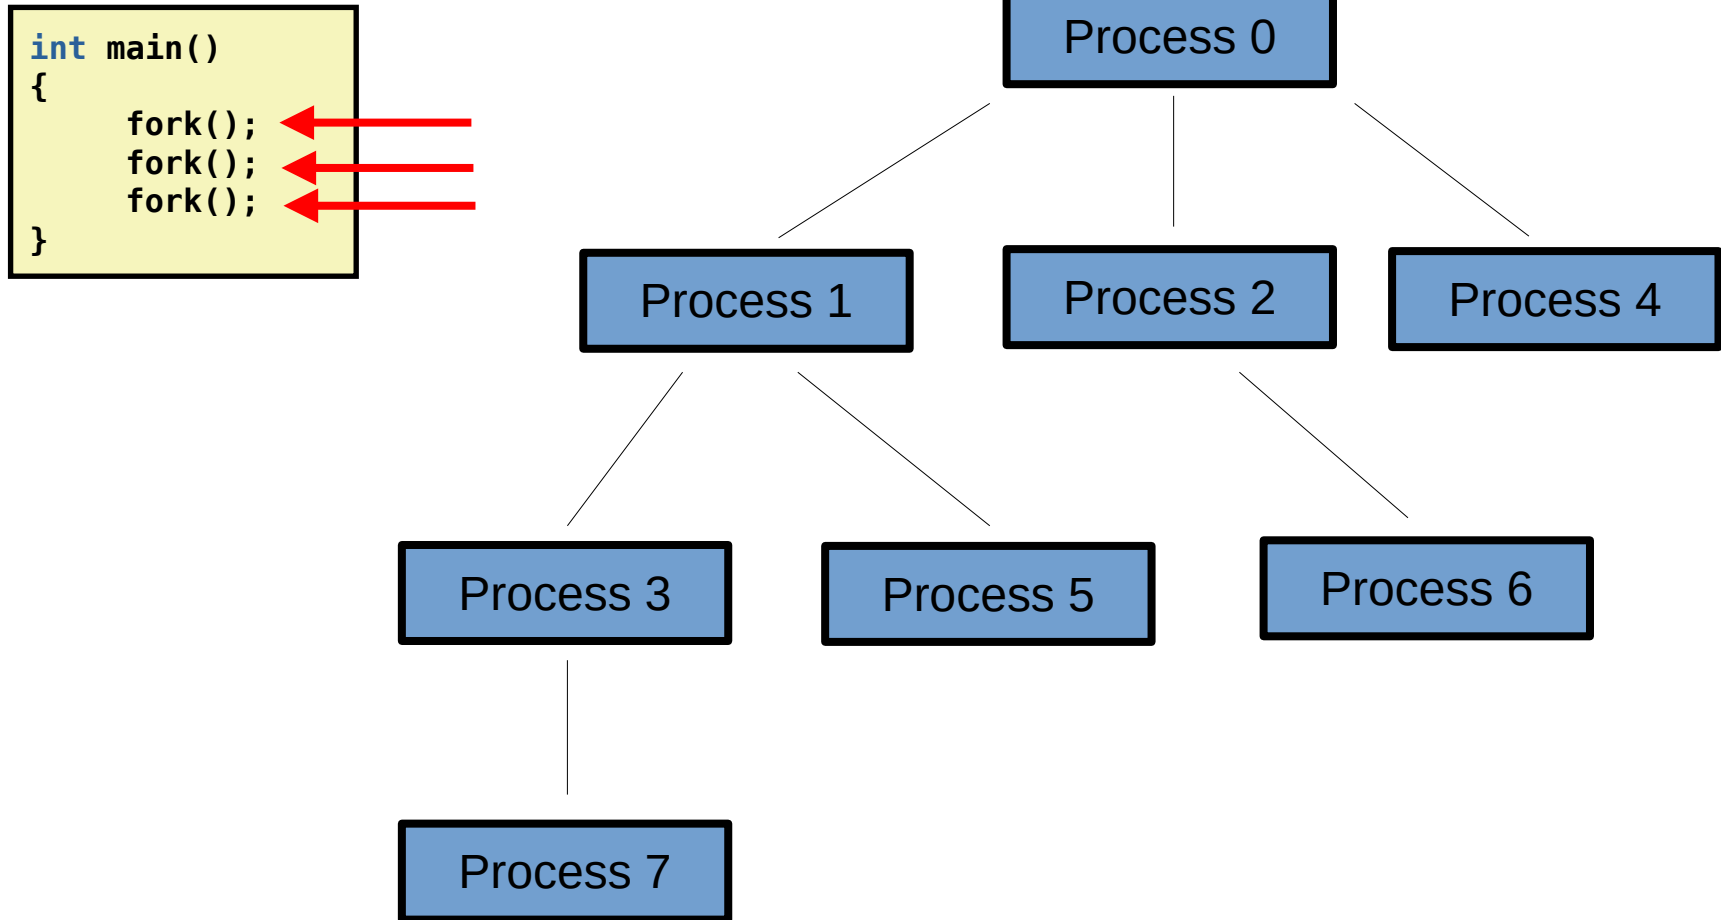
fork.c

```
user@system:~ gcc -o fork fork.c
user@system:~ ./fork
Hello from process 6196
Hello from process 6198
Hello from process 6197
Hello from process 6199
user@system:~
```

Processes in C

- **Very easy to get many different processes running concurrently (and in parallel if you've got the hardware)**
- **Like using OpenMP for threading, this can mean very few new lines of code**
- **Can be relatively easy to lose track of how many processes you've created**
- **Every `fork()` will create a child process**

Processes in C



Processes in C

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main()
{
    if (fork() == 0)
    {
        printf("This is the parent process");
    }
    fork();
    printf("Hello from process %d\n", getpid());
    exit(0);
}
```

parent.c

```
user@system:~ gcc -o parent parent.c
```

```
user@system:~ ./parent
```

```
Hello from process 13227
```

```
This is the parent process
```

```
Hello from process 13229
```

```
Hello from process 13228
```

```
Hello from process 13230
```

```
user@system:~
```

We don't want just copies

- **These models are good and all, but aren't always suitable**
 - Clients and Servers
 - Hardware simulations
 - Pipelining
- **We need a way of designing completely different processes from the ground up**
- **We could also run very different threads, but typically processes are used for this due to their differing nature**
- **C is not great at this so lets switch to Python (but it is perfectly possible in C!)**

Multiprocessing

- A Python library for creating multiple processes.
- Can be used to create pools of worker processes

```
import multiprocessing
import time

data = (
    ['A', '2'], ['B', '1'],
    ['C', '3'], ['D', '2']
)

def mp_worker(args):
    print(f"{args[0]} Waiting for {args[1]}s")
    time.sleep(int(args[1]))
    print(f"{args[0]} DONE")

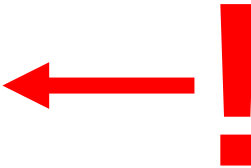
p = multiprocessing.Pool(2)
p.map(mp_worker, data)
```

pool.py

```
user@system:~ python3 pool.py
A Waiting for 2s
B Waiting for 1s
B DONE
C Waiting for 3s
A DONE
D Waiting for 2s
D DONE
C DONE
user@system:~
```

- But also can be used to define many different processes...

But we have a problem . . .

- Unlike threads, processes do not (in theory) share any data
- This means race conditions aren't a problem ← 
- But this does mean that sharing our data is going to be difficult
- **2 Strategies to deal with this:**
 - For child processes we can share channels / queues / pipes
 - For independent processes we can use sockets

Channels

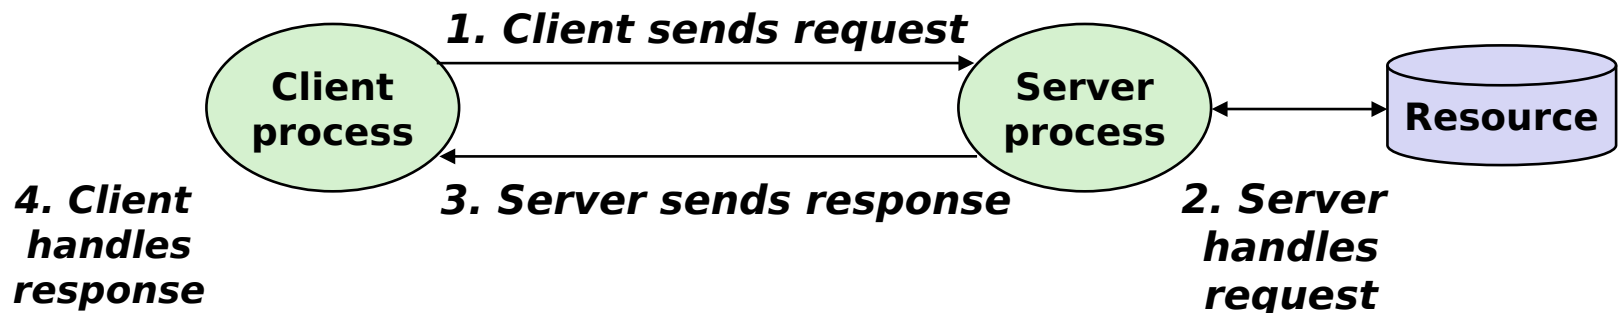
- Channels are the generic term we will use to describe any means for two processes to communicate
- Channels are (generally) one way communication tools



- Note these can also be used to eliminate race conditions if used instead of shared data (but with a lot of overhead)


A Client-Server Transaction

- **Most network applications are based on the client-server model:**
 - A **server** process and one or more **client** processes
 - Server manages some **resource**
 - Server provides **service** by manipulating resource for clients
 - Server activated by request from client (vending machine analogy)



Note: clients and servers are processes running on hosts (can be the same or different hosts)

A Client-Server Transaction



- Here, we are still going to use the client/server model
- Clients send messages, Servers receive them and reply
- Setting different processes up as separate programs can be handy, but sometimes we might wish to use several in one

A Multiprocessing Client-Server

```
import multiprocessing

PRODUCTION_COUNT = 4

def producer(to_consumer):
    for i in range(PRODUCTION_COUNT):
        to_consumer.put(f"Message {i}")

def consumer(from_producer):
    while True:
        message = from_producer.get()
        print(message)

q = multiprocessing.Queue()
process_list = [
    multiprocessing.Process(target=producer, args=(q,)),
    multiprocessing.Process(target=consumer, args=(q,))
]

for p in process_list:
    p.start()
```

```
user@system:~ python3 prod-cons.py
Message 0
Message 1
Message 2
Message 3
```

prod-cons.py

Pipes and Queues

- **Two semantically similar constructions for sending messages**
- **Both are essentially lists that can be added to and removed from**
- **Queues are much more lightweight but uni-directional**
- **Pipes are more computationally heavy, but bi-directional**
- **As we will see in a minute, bi-directionality is actually a problem if used carelessly**

Ending the Client-Server

```
import multiprocessing

PRODUCTION_COUNT = 4
KILL = "kill"

def producer(to_consumer):
    for i in range(PRODUCTION_COUNT):
        to_consumer.put(f"Message {I}")
    to_consumer.put(KILL)

def consumer(from_producer):
    while True:
        message = from_producer.get()
        if message == KILL:
            return
        print(message)


q = multiprocessing.Queue()
process_list = [
    multiprocessing.Process(target=producer, args=(q,)),
    multiprocessing.Process(target=consumer, args=(q,))
]

for p in process_list:
    p.start()
```

```
user@system:~ python3 prod-cons.py
Message 0
Message 1
Message 2
Message 3
user@system:~
```

prod-cons.py

A Process Communications



- Our example only used 2 processes, but we can use as many as we like and connect them however we like
- Of course nothing is ever that easy, we might run into problems if we do this without care

Blocking points


- **Process communication is blocking, that is sequential code will wait until both the client and the server are ready to proceed before it does so.**

Block points

```
...  
  
def producer(to_consumer):  
    for i in range(PRODUCTION_COUNT):  
        to_consumer.put(f"Message {I}")  
        to_consumer.put(KILL)  
  
def consumer(from_producer):  
    while True:  
        message = from_producer.get()  
        if message == KILL:  
            return  
        print(message)  
  
q = multiprocessing.Queue(1)  
  
...
```

Blocking points

- When drawing processes each blocking communication is marked as a channel either leaving or entering the processes
- Outbound communications leave the process
- Inbound communications enter the process
- Note that communications with replies are *usually* not shown separately



Blocking points

- A single process may have multiple blocking points
- These may each be to/from the same address, or to separate addresses
- They may be any combination of client / server communications


Block points

```
...
def producer(to_consumer):
    for i in range(PRODUCTION_COUNT):
        to_consumer.put(f"Message {I}")
        to_consumer.put(KILL)

def consumer(from_producer):
    while True:
        message = from_producer.get()
        if message == KILL:
            return
        print(message)
...
```

Blocking points

- A single process may have multiple blocking points
- These may each be to/from the same address, or to separate addresses
- They may be any combination of client / server communications




CSP

- **Communicating Sequential Processes**
- **First proposed by Tony Hoare in 1978**
- **Used as foundation for concurrency in many high level languages, such as Go**
- **Implementations in a variety of languages, but mostly outdated, so we won't use it directly**

- **No shared data**
- **Processes communicate with each other via channels**
- **If we map processes communications we can guarantee deadlock free**

CSP




CSP


- If we can map all communications, then we can ensure our design is sound
- If we never have a loop of Clients/Servers, we cannot deadlock
- This is mathmatically certain according to CSP
- Remember, we **must** assume that if deadlock might occur, it will

CSP


No Deadlock




Deadlock



No Deadlock




Deadlock



Avoiding Deadlock

- This harsh interpretation is not strictly true (but it is)
- For instance, networking has buffers and additional (semi-hidden) processes
- These are more visible in our multiprocessing example from earlier
- Queues, pipes, and networks almost always act as buffers allowing the client to progress



- This does not alter the conclusion though, its just slower to reach it

Avoiding Deadlock

- **Also, these diagrams do not actually mean we will deadlock, just that we might**
- **If we can avoid deadlock according to the diagram, we know we are deadlock free, guaranteed**
- **If not deadlock free according to the diagram, we just need to justify how we have avoided**
- **But the road to deadlock is paved with good intentions**

Ending the Client-Server

```
import multiprocessing
```

```
PRODUCTION_COUNT = 4
```

```
KILL = "kill"
```

```
def producer(to_consumer):  
    for i in range(PRODUCTION_COUNT):  
        to_consumer.put(f"Message {I}")  
    to_consumer.put(KILL)
```

```
def consumer(from_producer):  
    while True:  
        message = from_producer.get()  
        if message == KILL:  
            return  
        print(message)
```

```
q = multiprocessing.Queue()  
process_list = [  
    multiprocessing.Process(target=producer, args=(q,)),  
    multiprocessing.Process(target=consumer, args=(q,))  
]
```

```
for p in process_list:  
    p.start()
```

```
user@system:~ python3 prod-cons.py  
Message 0  
Message 1  
Message 2  
Message 3  
user@system:~
```

**This means we
don't really 'block'**

prod-cons.py

Broadcasting

- **Most communication channels are point to point**
- **They can be shared by many processes (as in multiple servers may share the same queue)**
- **But only a single processes will pull a single message.**
- **If you want a message to be received by multiple processes, it must be sent multiple times.**
- **Some libraries or channel types will allow a broadcast (one message to multiple servers), but this is not always possible**

Broadcasting

```
import multiprocessing
PRODUCTION_COUNT = 4
KILL = "kill"

def producer(to_consumer):
    for i in range(PRODUCTION_COUNT):
        to_consumer.put(f"Message {I}")
    to_consumer.put(KILL)
    to_consumer.put(KILL)
```

```
def consumer(from_producer):
    while True:
        message = from_producer.get()
        if message == KILL:
            print(f"{name} killed")
            return
        print(f"{name}: {message}")
```

```
q = multiprocessing.Queue()
process_list = [
    multiprocessing.Process(target=producer, args=(q,)),
    multiprocessing.Process(target=consumer, args=("A", q)),
    multiprocessing.Process(target=consumer, args=("B", q))
]
```

```
for p in process_list:
    p.start()
```



```
user@system:~ python3 mult-cons.py
A Message 0
A Message 1
B Message 2
A Message 3
B killed
A killed
user@system:~
```

Note that on a shared channel, there is no way to specifically address either consumer

mult-cons.py

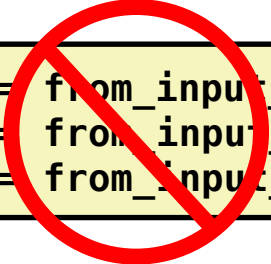
Broadcasting, or not

- However the lack of broadcasting can be a useful feature
- Only non-blocking (e.g. free processes) will pick up the message
- A useful property for if you want a thread-pool-like system



Choice

- **As well as multiple servers for a single client, we also have the case of multiple clients for one server**



```
message_1 = from_input_one.get()  
message_2 = from_input_two.get()  
message_3 = from_input_three.get()
```

- **We cannot try to read from each in turn, as this is a blocking operation**
- **But we will sometimes need to decide between multiple input channels**


Choice

- We can use **Select** statements to block until one or more channels are ready to read
- In Python this is some quite dense code

```
import select

read_list = [ ... ]
write_list = [ ... ]
exception_list = [ ... ]

read_ready, write_ready, exception_ready = select.select(
    read_list, write_list, exception_list)
```



Choice

```
import multiprocessing
import select

def producer(name, to_consumer):
    for i in range(3):
        to_consumer.put(f"{name} {I}")

def consumer(in_1, in_2):
    while True:
        (inputs, _, _) = select.select([in_1._reader, in_2._reader], [], [])
        if in_1._reader in inputs:
            message = in_1.get()
        elif in_2._reader in inputs:
            message = in_2.get()
        print(f"{message}")

q1 = multiprocessing.Queue()
q2 = multiprocessing.Queue()
process_list = [
    multiprocessing.Process(target=producer, args=("A", q1)),
    multiprocessing.Process(target=producer, args=("B", q2)),
    multiprocessing.Process(target=consumer, args=(q1, q2))
]

for p in process_list:
    p.start()
```

user@system:~ python3 select.py

A 0

A 1

B 0

A 2

B 1

B 2

select.py

Barriers

- **As well as just receiving isolated messages from multiple sources, it is often that we want to synchronise on multiple inputs**
- **For instance, how do we check that several processes are done**
- **We build a barrier and synchronise on that**

Barriers

```
import multiprocessing
import select
import time

def producer(name, sleepy_time, to_consumer):
    print(f"{name} sleep for {sleepy_time}")
    time.sleep(sleepy_time)
    print(f"{name} awoken")
    to_consumer.put(1)

def consumer(in_1, in_2):
    barrier[False, False]
    while True:
        (inputs, _, _) = select.select([in_1._reader, in_2._reader], [], [])
        if in_1._reader in inputs:
            _ = in_1.get()
            barrier[0] = True
        elif in_2._reader in inputs:
            _ = in_2.get()
            barrier[1] = True
        if all(i for i in barrier):
            print("Barrier passed")
            return

q1 = multiprocessing.Queue()
q2 = multiprocessing.Queue()
process_list = [
    multiprocessing.Process(target=producer, args=("A", 1, q1)),
    multiprocessing.Process(target=producer, args=("B", 2, q2)),
    multiprocessing.Process(target=consumer, args=(q1, q2))
]

for p in process_list:
    p.start()
```

```
user@system:~ python3 barrier.py
A sleep for 1
B sleep for 2
A Awoken
B Awoken
Barrier passed
user@system:~
```

Where would we use CSP-like systems?

Networking

- **Hopefully we've been through this enough already . . .**
- **All networked applications use this principle already**
- **Often time higher level network communications such as we used in A4 adds layers of complexity that hide these underlying principles but they're still there**
- **Small scale IOT devices and the like won't have enough resources to abstract them away though, and so they will be central**

Simulating Hardware

- **Hardware does not context switch, each component runs both concurrently and in parallel**
- **Therefore a good simulation of this would do the same**
- **These is a very common bachelors/masters projects**
 - Detector simulators (X-Rays, microscopes etc)
 - FPGA systems, processors, GPUs
 - Simulations of pre-production machines/experiments/products

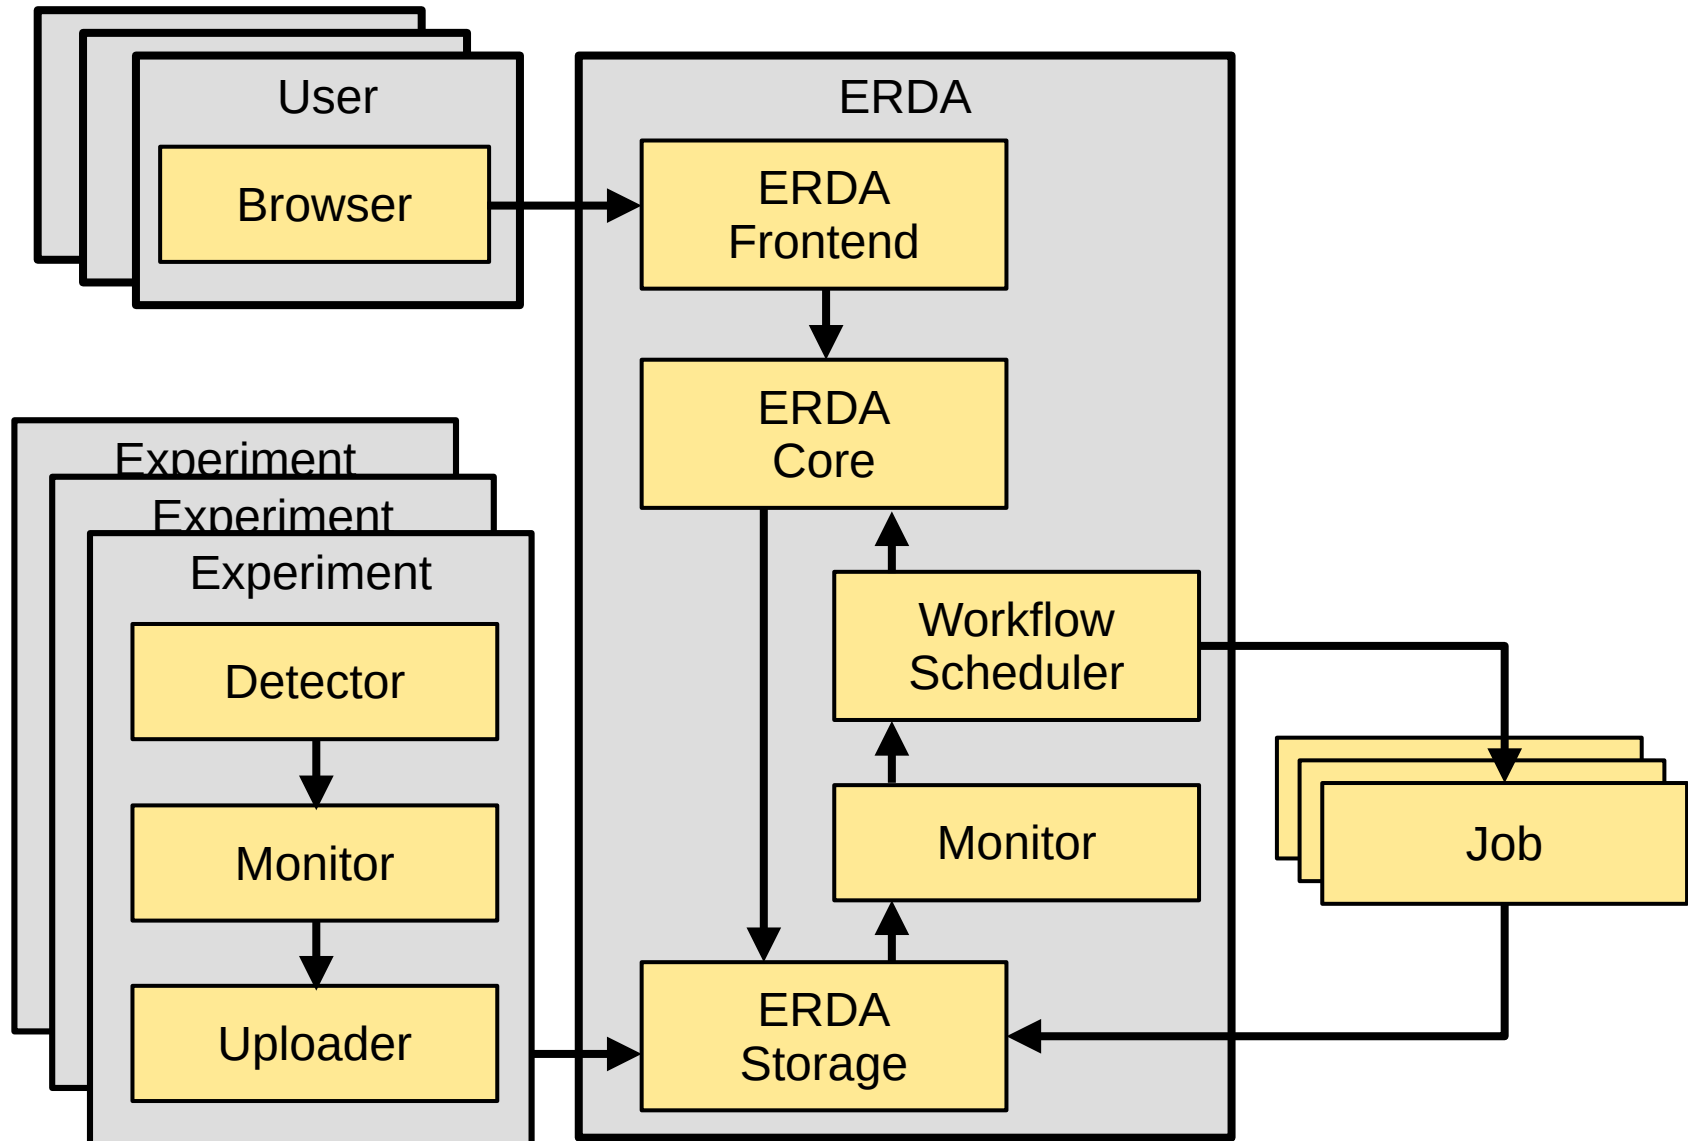
Pipelined Workflows

- **Scientific analysis can often be computed in a isolated, but dependent manner**
- **These is also a very common use case across all of science**
 - Simulation of physical systems (weather, astronomy)
 - Ongoing analysis
- **Also common in commercial space**
 - Big Data Analysis
 - Social Media
 - Stock Market Analysis

An Example Design

- **We won't look at the code here, it is long, complex, and dull**
- **This system encompasses many different machines and processes, communicating in a variety of ways**
- **Used to gather data, analyse the results dynamically and on an ongoing basis**
- **This system exists in its entirety, but as isolated parts and is being brought together to demonstrate the capability of this design methodology**

An Example Design



Conclusions

- **Going beyond simple loops introduces complexity**
- **OpenMP supports variation in looping**
 - Thread IDs can be used to identify individual jobs, elements, tasks, indexes etc etc etc
 - Chunking work is often a very good approach
 - Data parallel!
- **Inter process/thread communication is an alternative**
 - Much more complexity
 - Risks deadlock
 - But can be spread over different nodes/network
 - Task parallel!